

# SQL:

# Data Manipulation Language

Author: Diane Horton



UNIVERSITY OF  
TORONTO

[Slides 3-24 are covered by Prep4]

# Introduction

- So far, we have defined database schemas and queries mathematically.
- SQL is a formal language for doing so with a DBMS.
- “Structured Query Language”, but it’s for more than writing queries.
- Two sub-parts:
  - DDL (Data Definition Language), for defining schemas.
  - DML (Data Manipulation Language), for writing queries and modifying the database.

# PostgreSQL

- We'll be working in PostgreSQL, an open-source relational DBMS.
- Learn your way around the documentation; it will be very helpful.
- Standards?
  - There are several, the most recent being SQL:2008.
  - The standards are not freely available. Must purchase from the International Standards Organization (ISO).
  - PostgreSQL supports most of it SQL:2008.
  - DBMSs vary in the details around the edges, making portability difficult.

# Heads up: SELECT vs $\sigma$

- In SQL,
  - “SELECT” is for choosing columns, *i.e.*,  $\Pi$ .
  - Example:

```
SELECT surName
FROM Student
WHERE campus = 'StG';
```
- In relational algebra,
  - “select” means choosing rows, *i.e.*,  $\sigma$ .

# A high-level language

- SQL is a very high-level language.
  - Say “what” rather than “how.”
- You write queries without manipulating data. Contrast languages like Java or C++.
- Provides physical “data independence:”
  - Details of how the data is stored can change with no impact on your queries.
- You can focus on readability.
  - But because the DBMS optimizes your query, you get efficiency.

# Basic queries

# Meaning of a query with one relation

```
SELECT name  
FROM Course  
WHERE dept = 'CSC';
```

$\pi_{\text{name}} (\sigma_{\text{dept}=\text{"csc"}} (\text{Course}))$



## ... and with multiple relations

```
SELECT name  
FROM Offering, Took  
WHERE Offering.id = Took.oid and  
      dept = 'CSC';
```

$\Pi_{\text{name}} (\sigma_{\text{Offering.id=Took.id} \wedge \text{dept='csc'}} (\text{Offering} \times \text{Took}))$

# Temporarily renaming a table

- You can rename tables (just for the duration of the statement):

```
SELECT e.name, d.name  
FROM employee e, department d  
WHERE d.name = 'marketing'  
AND e.name = 'Horton';
```

- Can be convenient vs the longer full names:

```
SELECT employee.name, department.name  
FROM employee, department  
WHERE department.name = 'marketing'  
AND employee.name = 'Horton';
```

- This is like  $\rho$  in relational algebra.

# Self-joins

- As we know, renaming is *required* for self-joins.

- **Example:**

```
SELECT e1.name, e2.name  
FROM employee e1, employee e2  
WHERE e1.salary < e2.salary;
```

## \* In SELECT clauses

- A \* in the SELECT clause means “all attributes of this relation.”

- Example:

```
SELECT *  
FROM Course  
WHERE dept = 'CSC';
```

# Renaming attributes

- Use `AS «new name»` to rename an attribute in the result.

- **Example:**

```
SELECT name AS title, dept  
FROM Course  
WHERE breadth;
```

# Complex Conditions in a WHERE

- We can build boolean expressions with operators that produce boolean results.
  - comparison operators: `=`, `<>`, `<`, `>`, `<=`, `>=`
  - and many other operators:  
see section 6.1.2 of the text and chapter 9 of the postgresSQL documentation.
- Note that “not equals” is unusual: `<>`
- We can combine boolean expressions with:
  - Boolean operators: `AND`, `OR`, `NOT`.

# Example: Compound condition

- Find 3rd- and 4th-year CSC courses:

```
SELECT *  
FROM Offering  
WHERE dept = 'CSC' AND cnum >= 300;
```

# ORDER BY

- To put the tuples in order, add this as the final clause:

```
ORDER BY «attribute list» [DESC]
```

- The default is ascending order; DESC overrides it to force descending order.
- The attribute list can include expressions: e.g.,  

```
ORDER BY sales+rentals
```
- The ordering is the last thing done before the SELECT, so all attributes are still available.



# Case-sensitivity and whitespace

- Example query:

```
SELECT surName  
FROM Student  
WHERE campus = 'StG';
```

- Keywords, like `SELECT`, are not case-sensitive.
  - One convention is to use uppercase for keywords.
- Identifiers, like `Student` are not case-sensitive either.
  - One convention is to use lowercase for attributes, and a leading capital letter followed by lowercase for relations.
- Literal strings, like `'StG'`, are case-sensitive, and require single quotes.
- Whitespace (other than inside quotes) is ignored.

# Expressions in SELECT clauses

- Instead of a simple attribute name, you can use an expression in a SELECT clause.
- Operands: attributes, constants  
Operators: arithmetic ops, string ops

- **Examples:**

```
SELECT sid, grade+10 AS adjusted  
FROM Took;
```

```
SELECT dept || cnum  
FROM course;
```

# Expressions that are a constant

- Sometimes it makes sense for the whole expression to be a constant (something that doesn't involve any attributes!).
- **Example:**  

```
SELECT dept, cNum,  
       'satisfies' AS breadthRequirement  
FROM Course  
WHERE breadth;
```

# Pattern operators

- Two ways to compare a string to a pattern by:
  - `«attribute» LIKE «pattern»`
  - `«attribute» NOT LIKE «pattern»`
- Pattern is a quoted string
  - `%` means: any string
  - `_` means: any single character
- Example:

```
SELECT *  
FROM Course  
WHERE name LIKE '%Comp%';
```

# Aggregation and Grouping

# Computing on a column

- We often want to compute something across the values in a column.
- `SUM`, `AVG`, `COUNT`, `MIN`, and `MAX` can be applied to a column in a `SELECT` clause.
- Also, `COUNT ( * )` counts the number of tuples.
- We call this aggregation.
- Note: To stop duplicates from contributing to the aggregation, use `DISTINCT` inside the brackets. (Does not affect `MIN` or `MAX`.)

 **Example:** aggregation.txt

# Grouping

- **Example:** group-by.txt
- If we follow a SELECT-FROM-WHERE expression with GROUP BY <attributes>
  - The tuples are grouped according to the values of those attributes, and
  - any aggregation gives us a single value per group.

# Restrictions on aggregation

- If any aggregation is used, then each element of the **SELECT** list must be either:
  - aggregated, or
  - an attribute on the **GROUP BY** list.
- Otherwise, it doesn't even make sense to include the attribute.



# Continuing on with SQL

# HAVING Clauses

- **Example:** having.txt
- WHERE let's you decide which tuples to keep.
- Similarly, you can decide which *groups* to keep.
- Syntax:

...

GROUP BY «*attributes*»

HAVING «*condition*»

- Semantics:  
Only groups satisfying the condition are kept.

# Restrictions on HAVING clauses

- Outside subqueries, HAVING may refer to attributes only if they are either:
  - aggregated, or
  - an attribute on the GROUP BY list.
- (Same requirement as for SELECT clauses with aggregation)