

Subqueries

Where can a subquery go?

- Relational algebra syntax is so elegant that it's easy to see where subqueries can go.
- In SQL, a bit more thought is required . . .

Subqueries in a FROM clause

- In place of a relation name in the FROM clause, we can use a subquery.
- The subquery must be parenthesized.
- Must name the result, so you can refer to it in the outer query.

Worksheet, Q1:

```
SELECT sid, dept||cnum as course, grade
FROM Took,
    (SELECT *
     FROM Offering
     WHERE instructor='Horton') Hofferings
WHERE Took.oid = Hofferings.oid;
```

- This FROM is analogous to:
 $\text{Took} \times \rho_{\text{Hofferings}} (\llbracket \text{subquery} \rrbracket)$
- Can you suggest another version?

Subquery as a value in a WHERE

- If a subquery is guaranteed to produce exactly one tuple, then the subquery can be used as a value.
- Simplest situation: that one tuple has only one component.

Worksheet, Q2:

```
SELECT sid, surname
FROM Student
WHERE cgpa >
      (SELECT cgpa
       FROM Student
       WHERE sid = 99999);
```

- We can't do the analogous thing in RA:

$\Pi_{\text{sid, surname}} \sigma_{\text{cgpa} > (\text{«subquery»}) \text{Student}$

Special cases

- What if the subquery returns `NULL`?
- What if the subquery could return more than one value?

Quantifying over multiple results

- When a subquery can return multiple values, we can make comparisons using a quantifier.

- **Example:**

```
SELECT sid, surname
FROM Student
WHERE cgpa >
      (SELECT cgpa
       FROM Student
       WHERE campus = 'StG' );
```

- We can require that
 - $cgpa > \text{all of them}$, or
 - $cgpa > \text{at least one of them}$.

The Operator ANY

- Syntax:

$x \text{ «comparison» ANY («subquery»)$

or equivalently

$x \text{ «comparison» SOME («subquery»)$

- Semantics:

Its value is true iff the comparison holds for at least one tuple in the subquery result, i.e.,

$\exists y \in \text{«subquery results»} \mid x \text{ «comparison» } y$

- x can be a *list* of attributes,
but this feature is not supported by psql.

The Operator ALL

- Syntax:

$x \text{ «comparison» ALL («subquery»)}$

- Semantics:

Its value is true iff the comparison holds for every tuple in the subquery result, i.e.,

$\forall y \in \text{«subquery results»} \mid x \text{ «comparison» } y$

- x can be a list of attributes,
but this feature is not supported by psql.

• Example: **any-all**

The Operator IN

- Syntax:
 $x \text{ IN } (\langle\langle \text{subquery} \rangle\rangle)$
- Semantics:
 Its value is true iff x is in the set of rows generated by the subquery.
- x can be a list of attributes, and psql does support this feature.

Worksheet, Q3:

```
SELECT sid, dept||cnum AS course, grade
FROM Took NATURAL JOIN Offering
WHERE
    grade >= 80 AND
    (cnum, dept) IN (
        SELECT cnum, dept
        FROM Took NATURAL JOIN Offering
            NATURAL JOIN Student
        WHERE surname = 'Lakemeyer');
```

Worksheet, Q4:

Suppose we have tables $R(a, b)$ and $S(b, c)$.

1. What does this query do?

```
SELECT a
FROM R
WHERE b IN (SELECT b FROM S);
```

2. Can we express this query without using IN?

The Operator EXISTS

- Syntax:
EXISTS («*subquery*»)
- Semantics:
Its value is true iff the subquery has at least one tuple.
- Read it as “exists a row in the subquery result”

Example: EXISTS

```
SELECT surname, cgpa
FROM Student
WHERE EXISTS (
    SELECT *
    FROM Took
    WHERE Student.sid = Took.sid and
           grade > 85 );
```

Worksheet, Q5:

```
SELECT instructor
FROM Offering Off1
WHERE NOT EXISTS (
    SELECT *
    FROM Offering
    WHERE
        oid <> Off1.oid AND
        instructor = Off1.instructor );
```


Worksheet, Q6:

```
SELECT DISTINCT oid
FROM Took
WHERE EXISTS (
    SELECT *
    FROM Took t, Offering o
    WHERE
        t.oid = o.oid AND
        t.oid <> Took.oid AND
        o.dept = 'CSC' AND
        took.sid = t.sid );
```

x «comparison» ALL («subquery»)

$\forall y \in \text{«subquery results»} \mid x \text{ «comparison» } y$

x «comparison» SOME («subquery»)

$\exists y \in \text{«subquery results»} \mid x \text{ «comparison» } y$

x IN («subquery»)

Same as $x = \text{SOME} (\text{«subquery»})$

x NOT IN («subquery»)

Same as $x \neq \text{ALL} (\text{«subquery»})$

} just for
convenience

EXISTS («subquery»)

 $\exists y \in \text{«subquery results»}$

Scope

- If a name might refer to more than one thing, use the most closely nested one.
 - If a subquery refers only to names defined inside it, it can be evaluated **once** and used repeatedly in the outer query.
 - If it refers to any name defined outside of itself, it must be evaluated **once for each tuple in the outer query**.
- These are called **correlated subqueries**.

Renaming can make scope explicit

```
SELECT instructor
FROM Offering Off1
WHERE NOT EXISTS (
    SELECT *
    FROM Offering Off2
    WHERE
        Off2.oid <> Off1.oid AND
        Off2.instructor = Off1.instructor );
```

Summary: where subqueries can go

- As a relation in a FROM clause.
- As a value in a WHERE clause.
- With ANY, ALL, IN or EXISTS in a WHERE clause.
- As operands to UNION, INTERSECT or EXCEPT.
- Reference: textbook, section 6.3.